

OpenMP Programming Assignment

Organization: CS312, Computer Science, SIUE
Author: Mark McKenney

Overview:

Write a multithreaded C++ program using open MP threads. I will provide a text file containing numbers that must be read into a two dimensional array. The first line of the file will have 2 integers, the number of rows in the array followed by the number of columns in the array. The rest of the file will contain integers that the array must be initialized to. Your program must also utilize a random number generator. I will provide source code.

Description:

Your program will indicate the cell address (the row and column) of the cell with the highest neighborhood average. If there is a tie, your program must report only a single cell that ties the maximum value. I will post a solution consisting of all cells that tie the maximum value so you can be sure your solution is correct. The neighborhood of a cell is all cells that immediately border the cell, including the cell itself. For each cell, you must compute the average of numbers in the neighborhood. For example, in the following array:

```
unsigned int M[10000][10000];  
the neighborhood of cell M[2][8] consists of the following cells:
```

```
M[1][7] M[1][8] M[1][9]  
M[2][7] M[2][8] M[2][9]  
M[3][7] M[3][8] M[3][9]
```

Because the array will be large, you will need to use dynamic memory allocation (keyword new) on the heap.

Be careful not to go out of bounds on the array when computing neighborhoods

Your program should take a command line argument indicating the number of threads that will be used.

The group that has the fastest program when run with the number of threads equal to the number of available cores on a chosen computer will win a prize (of non-monetary value)! So don't give away your speedy secrets.

The openMP wiki has a nice intro to using openMP. <http://en.wikipedia.org/wiki/OpenMP>

Note that you must have a GCC compiler version 4.3 or later to use openMP. The department's home.cs.siue.edu server has a sufficient GCC version. You

can download the newest version of GCC for mac and linux. You may have to install cygwin to do this with windows. Microsoft's compiler has it too.

Sample input files are attached

Your program will be graded on home.cs.siue.edu, so make sure it compiles and runs correctly there. Your group must make exactly 1 submission, and your group member names should be in a comment in the first line of your file. Also, you **MUST** submit a MAKEFILE. A sample makefile is provided, you can use this one, or edit it for your needs. Your submission should be zipped and submitted as single file.

What to Turn In:

You must turn in your source code, a Makefile that will compile your source code on home.cs.siue.edu, and a README file that contains the names of your group members, and any additional information about your implementation that you think I may need to consider while grading.

Turn in a single zip or tar.gz file.

The README file should be a PLAIN TEXT FILE!

Example Input and Expected Output:

Here are some sample runs of the program so you can see what the output looks like. Your output should look **EXACTLY** like mine, except that you will only show a single cell address with the largest average. The input file small3.txt contains two cells tying for the largest average

The following shows the expected output with the timing output. Your output should match these. If you can beat these times, you are doing very well!

Output that does match the format shown below will receive a grade of 0.

```
vm-02$ cat Makefile
```

```
all:
```

```
    g++ -fopenmp -ggdb  matAverager.cpp -o matavg
```

```
vm-02$ make
```

```
g++ -fopenmp -ggdb  matAverager.cpp -o matavg
```

```
vm-02$ ./matavg
```

```
usage: exe [input data file] [num of threads to use]
```

```
or usage: exe rand [num of threads to use] [num rows] [num cols] [seed value]
```

```
vm-02$ ./matavg small3.txt 10
```

```
largest average: 7.66667
```

```
found at cells: (0,1) (0,2)
```

```
elapsed time: 0.033653
```

```
vm-02$ ./matavg rand 10 5 5 0
```

```
largest average: 6780.75
```

```
found at cells: (4,4)
```

```
elapsed time: 0.0271368
```

```

vm-02$ ./matavg rand 10 100 100 0
largest average: 8296.33
found at cells: (99,2)
elapsed time: 0.0396051
vm-02$ ./matavg rand 10 1000 1000 0
largest average: 9042.44
found at cells: (966,225)
elapsed time: 0.08917
vm-02$ ./matavg rand 10 1000 2000 0
largest average: 9339.83
found at cells: (999,1504)
elapsed time: 0.080462
vm-02$ ./matavg rand 1 1000 2000 0
largest average: 9339.83
found at cells: (999,1504)
elapsed time: 0.162964
vm-02$ ./matavg rand 4 1000 2000 3
largest average: 9370.22
found at cells: (874,617)
elapsed time: 0.077152
vm-02$ ./matavg rand 1 1000 2000 3
largest average: 9370.22
found at cells: (874,617)
elapsed time: 0.155171
vm-02$ ./matavg rand 10 1000 2000 3
largest average: 9370.22
found at cells: (874,617)
elapsed time: 0.0878711
vm-02$ ./matavg rand 10 10000 2000 0
largest average: 9240.89
found at cells: (8524,739)
elapsed time: 0.336246
vm-02$ ./matavg rand 1 10000 2000 0
largest average: 9240.89
found at cells: (8524,739)
elapsed time: 2.0535
vm-02$ ./matavg rand 10 10000 20000 0
largest average: 9461.78
found at cells: (618,2726)
elapsed time: 2.48621
vm-02$

```

Additional Files for this Project:

```
#include <sys/time.h>
#include <fstream>
#include <iostream>
#include <omp.h>
#include <cstdlib>
#include <sstream>
#include <list>
#include <string>

using namespace std;

// a class to get more accurate time

class stopwatch{
private:
    double elapsedTime;
    double startedTime;
    bool timing;
    //returns current time in seconds
    double current_time( )
    {
        timeval tv;
        gettimeofday(&tv, NULL);
        double rtn_value = (double) tv.tv_usec;
        rtn_value /= 1e6;
        rtn_value += (double) tv.tv_sec;
        return rtn_value;
    }
public:
    stopwatch( ): elapsedTime( 0 ), startedTime( 0 ), timing( false )
    {
    }

    void start( )
    {
        if( !timing )
        {
            timing = true;
            startedTime = current_time( );
        }
    }

    void stop( )
    {
        if( timing )
        {
            elapsedTime += current_time( )-startedTime;
            timing = false;
        }
    }

    void resume( )
    {
        start( );
    }

    void reset( )
    {
        elapsedTime = 0;
        startedTime = 0;
        timing = false;
    }
}
```

```
double getTime( )
{
    return elapsedTime;
}

};

// function takes an array pointer, and the number of rows and cols in the array, and
// allocates and initializes the two dimensional array to a bunch of random numbers

void makeRandArray( unsigned int **& data, unsigned int rows, unsigned int cols,
unsigned int seed )
{
    // allocate the array
    data = new unsigned int*[ rows ];
    for( unsigned int i = 0; i < rows; i++ )
    {
        data[i] = new unsigned int[ cols ];
    }

    // seed the number generator
    // you should change the seed to get different values
    srand( seed );

    // populate the array
    for( unsigned int i = 0; i < rows; i++ )
        for( unsigned int j = 0; j < cols; j++ )
        {
            data[i][j] = rand() % 10000 + 1; // number between 1 and 10000
        }
}

void getDataFromFile( unsigned int **& data, char fileName[], unsigned int &rows,
unsigned int &cols )
{
    ifstream in;
    in.open( fileName );
    if( !in )
    {
        cerr << "error opening file: " << fileName << endl;
        exit( -1 );
    }

    in >> rows >> cols;
    data = new unsigned int*[ rows ];
    for( unsigned int i = 0; i < rows; i++ )
    {
        data[i] = new unsigned int[ cols ];
    }

    // now read in the data
    for( unsigned int i = 0; i < rows; i++ )
        for( unsigned int j = 0; j < cols; j++ )
        {
            in >> data[i][j];
        }
}
```

```

int main( int argc, char* argv[] )
{
    if( argc < 3 )
    {
        cerr<<"Usage: " << argv[0] << " [input data file] [num of threads to
use] " << endl;

        cerr<<"or" << endl << "Usage: " << argv[0] << " rand [num of threads to
use] [num rows] [num cols] [seed value]" << endl;
        exit( 0 );
    }

    // read in the file
    unsigned int rows, cols, seed;
    unsigned int numThreads;
    unsigned int ** data;
    // convert numThreads to int
    {
        stringstream ssl;
        ssl << argv[2];
        ssl >> numThreads;
    }

    string fName( argv[1] );
    if( fName == "rand" )
    {
        {
            stringstream ssl;
            ssl << argv[3];
            ssl >> rows;
        }
        {
            stringstream ssl;
            ssl << argv[4];
            ssl >> cols;
        }
        {
            stringstream ssl;
            ssl << argv[5];
            ssl >> seed;
        }
        makeRandArray( data, rows, cols, seed );
    }
    else
    {
        getDataFromFile( data, argv[1], rows, cols );
    }

    cerr << "data: " << endl;
    for( unsigned int i = 0; i < rows; i++ )
    {
        for( unsigned int j = 0; j < cols; j++ )
        {
            cerr << "i,j,data " << i << ", " << j << ", ";
            cerr << data[i][j] << " ";
        }
        cerr << endl;
    }
    cerr<< endl;

    // tell omp how many threads to use
    omp_set_num_threads( numThreads );

#pragma omp parallel
{

```

```
    }

    stopwatch S1;
    S1.start();

    //////////////////////////////////////
    ////////////////////////////////////// YOUR CODE HERE //////////////////////////////////////
    //////////////////////////////////////
    cout << "ASDFGASDFSADF"<<endl;

    S1.stop();

    // print out the max value here

    cerr << "elapsed time: " << S1.getTime( ) << endl;
}
```

```
all:
    g++ -fopenmp -ggdb matAverager.cpp -o matavg
```


small.txt

Page 1 of 1

```
4 4
4 9 7 4
8 9 5 9
4 2 5 5
9 6 8 2
```

small2.txt

Page 1 of 1

```
10 10
42876 80635 79090 55398 88954 60207 42109 78347 24948 83861
84807 83387 75906 77075 81374 78301 6499 46947 75625 41873
83385 40154 24212 89395 39855 91713 45647 21516 33935 87363
13206 28824 98607 52530 47248 2899 27395 6 70589 4602
65547 69509 56011 41134 87304 40822 80754 88828 55322 46952
24351 74639 4755 74920 52453 91664 18970 94612 71803 28076
93594 29326 67589 31367 55106 75102 26878 97867 30321 25078
21074 98872 65557 67195 17754 83799 67330 56550 94735 61454
93786 39378 73680 20705 85303 83209 78279 63671 54949 34690
49007 50233 64178 84813 53042 7961 7972 20015 28102 86594
```

small3.txt

Page 1 of 1

```
4 4
4 9 7 4
8 9 9 8
4 2 2 4
9 6 6 9
```

```
import sys
import random

if len( sys.argv ) >= 4:
    fileName = sys.argv[1]
    rows = int( sys.argv[2] )
    cols = int( sys.argv[3] )
else:
    print 'usage: exe [output filename] [num of rows] [num of cols]'
    exit()

file = open( fileName, 'w' )

random.seed( )

file.write( str( rows ) + ' ' )
file.write( str( cols ) + "\n" )

for i in range( 0,rows ):
    for j in range( 0, cols ):
        file.write( str( random.randrange(2,100000) ) )
        file.write( ' ' )
    file.write( "\n" )
```

timeDiff.bash

Page 1 of 1

```
#!/bin/bash
START=$(date +%s)
# do something

# start your script work here
$1 $2 $3 $4 $5 $6
# your logic ends here

END=$(date +%s)
DIFF=$(( $END - $START ))
echo "It took $DIFF seconds"
```